

IMPLEMENTATION OF A NEW SYSTEM CALL

Submitted By:

K.A.SIVA VARDHAN-B160333CS

B.MAHESWARARAO-B160349CS

R.VAMSI KRISHNA-B160109CS

P.SATYANARAYANA-B160340CS

INTRODUCTION

- Implementing a new system call called *sys_hello* .
- This system call helps us to know the FILE NAME and PID of the runnable, terminated, sleeping and blocked processes along with their total count . Along with this the user also get to know the information of the high priority , static priority , and normal priority of the processes along with their count.

STEPS INVOLVED

1. Creating a directory which contains source code.
2. Linking the new system call with kernel .
3. Modifying the file *syscall_64.tbl* .
4. Modifying the file *syscalls.h* .
5. Compiling and booting.

CREATING DIRECTORY

- Create a new directory “*hello*” and open that directory.
- Create a new file “*hello.c*” and write the source code for the system call in that.
- Create a Makefile which contains:
 obj -y := hello.o
- It implies that our code is compiled and included in the source code.

LINKING THE SYSTEM CALL TO KERNEL

- Add *hello* directory to the Makefile .
- Through this we are specifying to the compiler that the source files of system call are present in *hello* directory.

```
945     EXPORT_SYMBOL_STACK_VALIDATION
946 endif
947 endif
948
949
950 ifeq ($(KBUILD_EXTMOD),)
951 core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/
952
953 vmlinux-dirs := $(patsubst %/,%, $(filter %/, $(init-y) $(init-m) \
954     $(core-y) $(core-m) $(drivers-y) $(drivers-m) \
955     $(net-y) $(net-m) $(libs-y) $(libs-m) $(virt-y)))
956
957 vmlinux-alldirs := $(sort $(vmlinux-dirs) $(patsubst %/,%, $(filter %/, \
```

MODIFYING *SYSCALL_64.TBL*

- It is located in `/arch/x86/entry/syscalls` .
- You can also know it using the command:
“find -name syscall_64.tbl”
- Include the new system call number and it's entry point.

```

324 313 common rlimit_module      __x64_sys_rlimit_module
325 314 common sched_setattr      __x64_sys_sched_setattr
326 315 common sched_getattr      __x64_sys_sched_getattr
327 316 common renameat2          __x64_sys_renameat2
328 317 common seccomp            __x64_sys_seccomp
329 318 common getrandom           __x64_sys_getrandom
330 319 common memfd_create        __x64_sys_memfd_create
331 320 common kexec_file_load     __x64_sys_kexec_file_load
332 321 common bpf                __x64_sys_bpf
333 322 64  execveat              __x64_sys_execveat/ptregs
334 323 common userfaultfd        __x64_sys_userfaultfd
335 324 common membarrier          __x64_sys_membarrier
336 325 common mlock2             __x64_sys_mlock2
337 326 common copy_file_range    __x64_sys_copy_file_range
338 327 64  preadv2              __x64_sys_preadv2
339 328 64  pwritev2             __x64_sys_pwritev2
340 329 common pkey_mprotect      __x64_sys_pkey_mprotect
341 330 common pkey_alloc          __x64_sys_pkey_alloc
342 331 common pkey_free          __x64_sys_pkey_free
343 332 common statx              __x64_sys_statx
344 333 common io_pgetevents      __x64_sys_io_pgetevents
345 334 common rseq               __x64_sys_rseq
346 335 64  hello                sys_hello
347
348 #
349 # x32-specific system call numbers start at 512 to avoid cache impact
350 # for native 64-bit operation. The __x32_compat_sys stubs are created
351 # on-the-fly for compat_sys_*(()) compatibility system calls if X86_X32
352 # is defined.
353 #
354 512 x32 rt_sigaction          __x32_compat_sys_rt_sigaction
355 513 x32 rt_sigreturn          sys32_x32_rt_sigreturn
356 514 x32 ioctl                __x32_compat_sys_ioctl
357 515 x32 readv                __x32_compat_sys_readv
358 516 x32 writev               __x32_compat_sys_writev
359 517 x32 recvfrom              __x32_compat_sys_recvfrom

```

MODIFYING *SYSCALLS.H*

- It is located in `/include/linux` .
- Add the following line at the end of the file.
asm linkage long sys_hello(void) ;


```
asmlinkage long sys_old_mmap(struct mmap_arg_struct __user *arg);
```

```
/*  
 * Not a real system call, but a placeholder for syscalls which are  
 * not implemented -- see kernel/sys_ni.c  
 */
```

```
asmlinkage long sys_ni_syscall(void);
```

```
asmlinkage long sys_hello(void);
```

```
#endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */
```

```
/*  
 * Kernel code should not call syscalls (i.e., sys_xyzzyz()) directly.  
 * Instead, use one of the functions which work equivalently, such as  
 * the ksys_xyzzyz() functions prototyped below.  
 */
```

```
int ksys_mount(char __user *dev_name, char __user *dir_name, char __user *type,  
               unsigned long flags, void __user *data);
```

```
int ksys_umount(char __user *name, int flags):
```

COMPILING AND BOOTING

- To use the new system call we should recompile the kernel first.
- Use the following commands for that:
 - i. `sudo make -j 4`***
 - ii. `sudo make modules_install -j 4`***
 - iii. `sudo make install -j 4`***
 - iv. `sudo update-grub`***
- Restart the system.

TESTING OUR NEW SYSTEM CALL

- To test the new system call write a simple program “*test.c*”.
- Compile and execute this program.
- If it runs successfully, it’ll give the corresponding cause and we can use the ‘*dmesg*’ command to check the kernel log.

```
1 #include<stdio.h>
2 #include<linux/kernel.h>
3 #include<sys/syscall.h>
4 #include<unistd.h>
5 int main()
6 {
7     printf("Invoking 'listProcessInfo' system call\n");
8     long int ret_status=syscall(335);
9     if(ret_status==0)
10         printf("System call 'listProcessInfo' executed correctly. Use dmesg to check processInfo\n");
11     else
12         printf("System call 'listProcessInfo' did not execute as expected\n");
13     return 0;
14 }
```

THANK YOU

A series of horizontal lines in teal and light blue colors, with varying lengths and thicknesses, extending across the width of the slide below the text.